

P8106 - HW1

Ravi Brenner

Contents

Data and problem	1
a. Lasso model	2
b. Elastic net	5
c. Partial least squares	8
d. Model comparison	9
e. Lasso using glmnet	11

```
library(glmnet)
library(caret)
library(tidymodels)
library(pls)
library(readr)
```

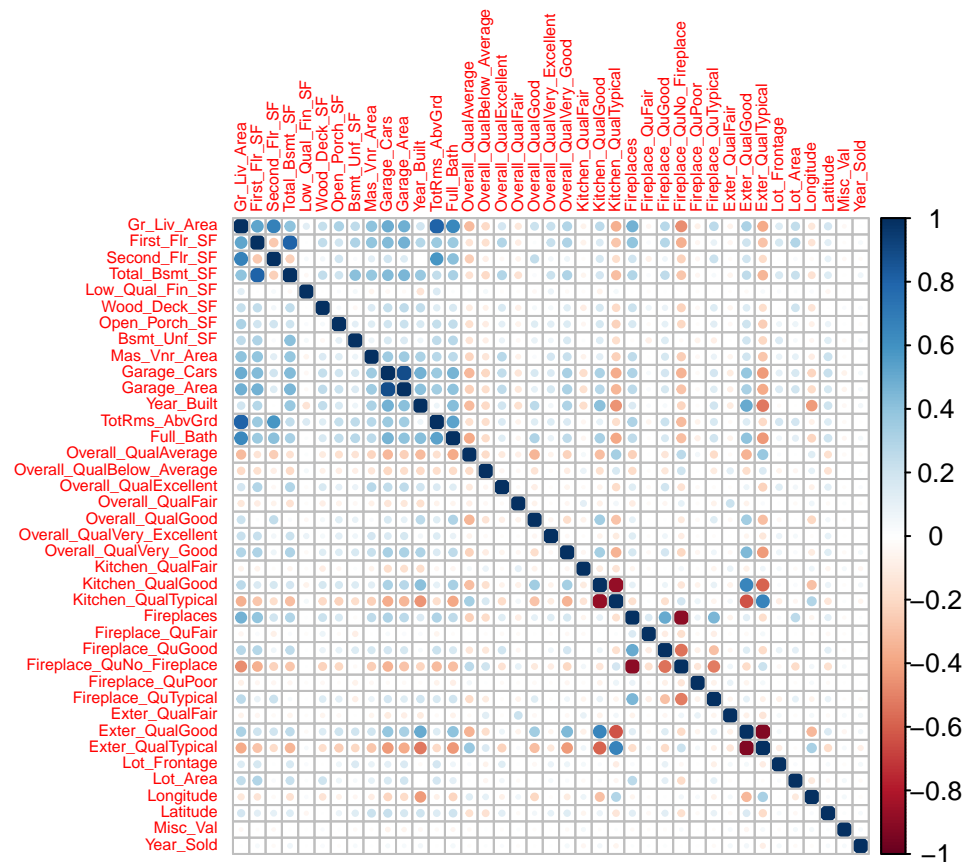
Data and problem

Load data

```
training <- read_csv("housing_training.csv")
testing <- read_csv("housing_test.csv")
```

Very brief exploration of the training data

```
x <- model.matrix(Sale_Price ~ ., training)[,-1]
corrplot::corrplot(cor(x),
                    method = "circle", type = "full", tl.cex = 0.5)
```



a. Lasso model

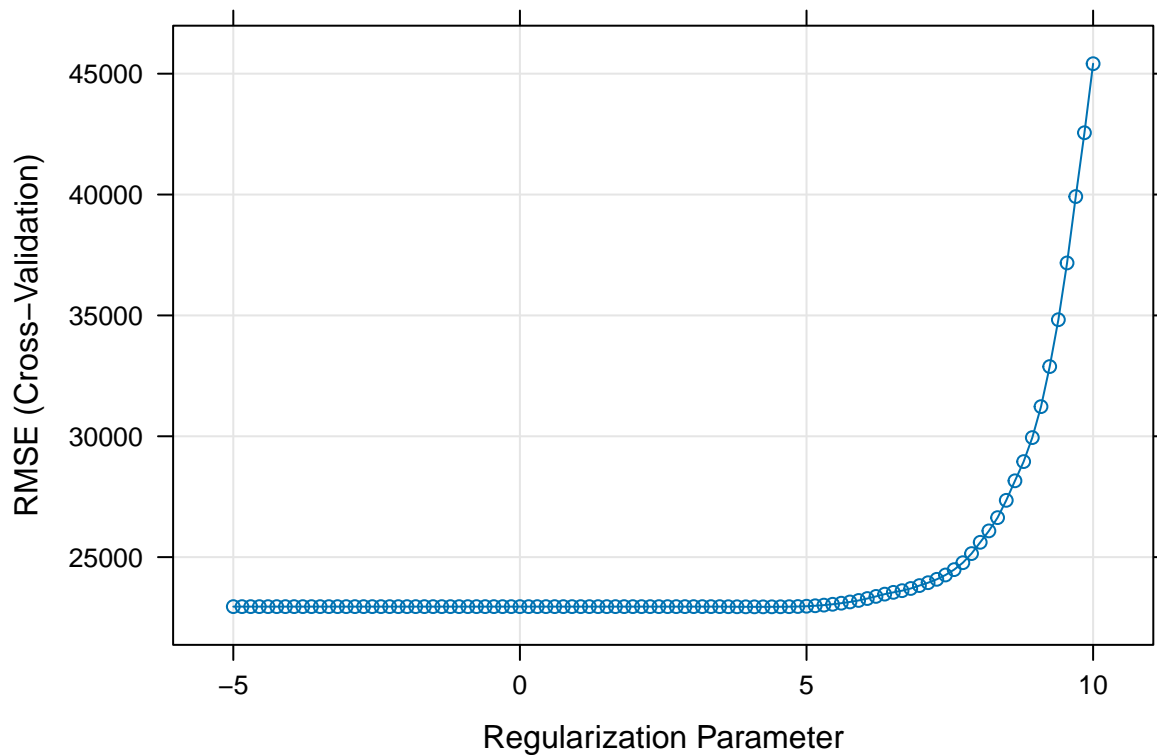
- (a) Fit a lasso model on the training data. Report the selected tuning parameter and the test error. When the 1SE rule is applied, how many predictors are included in the model?

Fit model

```
ctrl1 <- trainControl(method = "cv", number = 10, selectionFunction = "best")

set.seed(2025)
lasso.fit <- train(Sale_Price ~ .,
  data = training,
  method = "glmnet",
  tuneGrid = expand.grid(alpha = 1,
    lambda = exp(seq(10, -5, length = 100))),
  trControl = ctrl1)

plot(lasso.fit, xTrans = log)
```



Best tuning parameter

```
best_lambda <- lasso.fit$bestTune$lambda
best_lambda
```

```
## [1] 69.57632
```

test error

```
lasso.pred <- predict(lasso.fit, newdata = testing)
mean((lasso.pred - testing[, "Sale_Price"])^2 |> pull()) #MSE
```

```
## [1] 439568493
```

```
mean((lasso.pred - testing[, "Sale_Price"])^2 |> pull()) |> sqrt() #RMSE
```

```
## [1] 20965.89
```

How many coefficients in 1se model?

First, find 1se lambda value

```
# fit 1se version
ctrl2 <- trainControl(method = "cv", number = 10, selectionFunction = "oneSE")
set.seed(2025)
lasso.1se <- train(Sale_Price ~ .,
  data = training,
  method = "glmnet",
  tuneGrid = expand.grid(alpha = 1,
    lambda = exp(seq(10, -5, length = 100))),
  trControl = ctrl2)

lambda_1se <- lasso.1se$bestTune |> pull(lambda)
```

We can also calculate this value manually. The model output `results` gives the RMSE and RMSESD. To convert the SD to SE, we have to divide by the square root (because it is *root* mean square error) of the square root of the fold size (the standard way to calculate SE from SD). In this case the fold size is 144, so we divide by the square root of 12.

```
rmse_1se <- lasso.fit$results |>
  filter(lambda == best_lambda) |>
  mutate(RMSESE = RMSESD / sqrt(12),
         RMSE_1SE = RMSE + RMSESE) |>
  pull(RMSE_1SE)

lasso.fit$results |>
  filter(RMSE > 23500,
         RMSE < 23600)
```

	alpha	lambda	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
## 1	1	675.2963	23542.89	0.8973281	16778.67	2281.745	0.01771693	1527.013

Model with 1se lambda value:

```
coef(lasso.fit$finalModel, lambda_1se)

## 40 x 1 sparse Matrix of class "dgCMatrix"
##                               s1
## (Intercept)                -2.916836e+06
## Gr_Liv_Area                  5.794761e+01
## First_Flr_SF                 1.006061e+00
## Second_Flr_SF                .
## Total_Bsmt_SF                3.674362e+01
## Low_Qual_Fin_SF             -2.879529e+01
## Wood_Deck_SF                 8.702628e+00
## Open_Porch_SF               9.147142e+00
## Bsmt_Unf_SF                 -2.005078e+01
## Mas_Vnr_Area                 1.402927e+01
## Garage_Cars                  3.115356e+03
## Garage_Area                  1.082411e+01
## Year_Built                   3.110661e+02
## TotRms_AbvGrd               -1.695581e+03
## Full_Bath                    .
## Overall_QualAverage          -3.311972e+03
## Overall_QualBelow_Average    -9.646160e+03
## Overall_QualExcellent        9.132387e+04
## Overall_QualFair             -7.183517e+03
## Overall_QualGood              1.031218e+04
## Overall_QualVery_Excellent   1.614854e+05
## Overall_QualVery_Good        3.677171e+04
## Kitchen_QualFair             -5.938606e+03
## Kitchen_QualGood              .
## Kitchen_QualTypical          -9.520503e+03
## Fireplaces                   6.652772e+03
## Fireplace_QuFair             -1.054475e+02
## Fireplace_QuGood              4.425741e+03
## Fireplace_QuNo_Fireplace     .
## Fireplace_QuPoor              .
## Fireplace_QuTypical          -9.329979e+02
## Exter_QualFair               -1.522564e+04
```

```
## Exter_QualGood          .
## Exter_QualTypical      -4.883074e+03
## Lot_Frontage           7.603302e+01
## Lot_Area               5.739267e-01
## Longitude              -1.428573e+04
## Latitude               2.408926e+04
## Misc_Val               .
## Year_Sold               .

# Number of predictors
coef(lasso.fit$finalModel, lambda_1se) |> as.matrix() |> as.data.frame() |> filter(s1 > 0) |> nrow()

## [1] 18
```

There are 18 predictors in the 1SE lasso model, compared to 25 predictors in the whole data set.

b. Elastic net

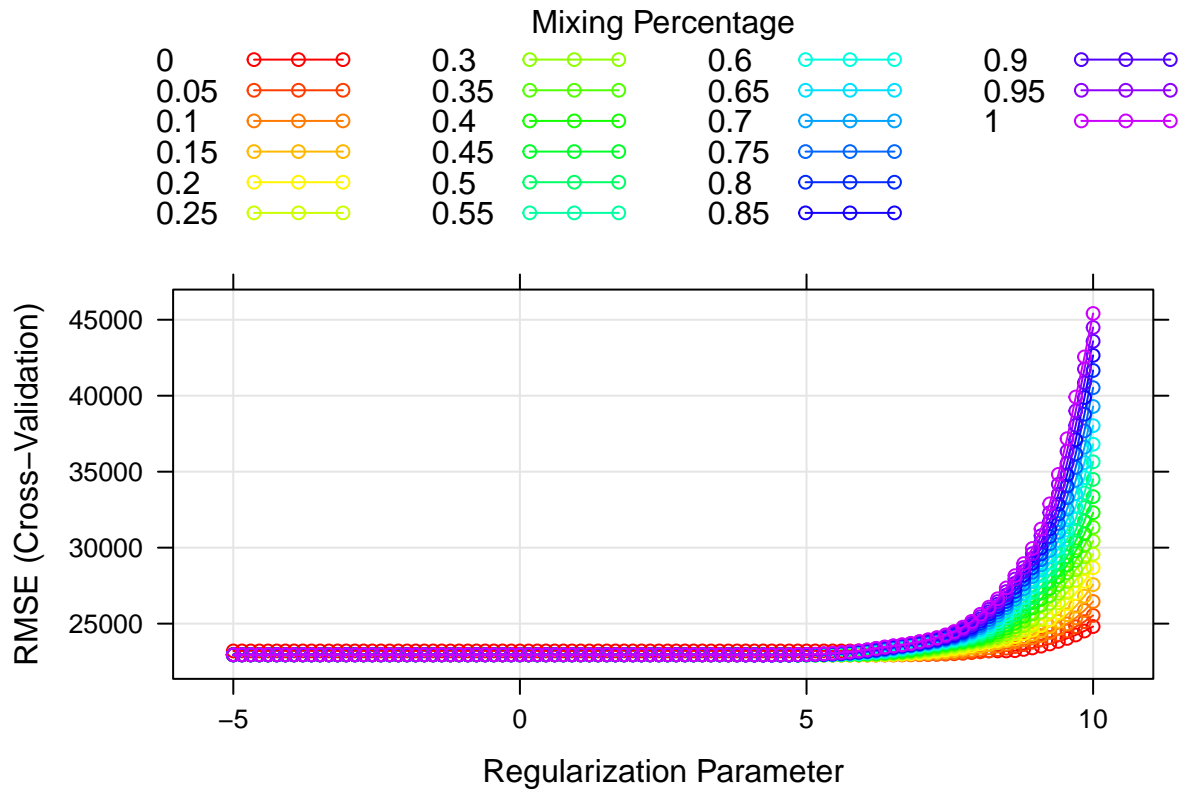
- (b) Fit an elastic net model on the training data. Report the selected tuning parameters and the test error. Is it possible to apply the 1SE rule to select the tuning parameters for elastic net? If the 1SE rule is applicable, implement it to select the tuning parameters. If not, explain why.

Fit Elastic net model

```
set.seed(2025)
enet.fit <- train(Sale_Price ~ .,
  data = training,
  method = "glmnet",
  tuneGrid = expand.grid(alpha = seq(0, 1, length = 21),
    lambda = exp(seq(10, -5, length = 100))),
  trControl = ctrl1)

myCol <- rainbow(25)
myPar <- list(superpose.symbol = list(col = myCol),
  superpose.line = list(col = myCol))

plot(enet.fit, par.settings = myPar, xTrans = log)
```



```
# coefficients in the final model
coef(enet.fit$finalModel, s = enet.fit$bestTune$lambda)
```

```
## 40 x 1 sparse Matrix of class "dgCMatrix"
##                               s1
## (Intercept)                -5.112158e+06
## Gr_Liv_Area                  3.877699e+01
## First_Flr_SF                 2.670023e+01
## Second_Flr_SF               2.545573e+01
## Total_Bsmt_SF               3.494355e+01
## Low_Qual_Fin_SF             -1.586366e+01
## Wood_Deck_SF                1.232450e+01
## Open_Porch_SF               1.687896e+01
## Bsmt_Unf_SF                 -2.072832e+01
## Mas_Vnr_Area                1.165193e+01
## Garage_Cars                 4.046992e+03
## Garage_Area                 8.889587e+00
## Year_Built                  3.192146e+02
## TotRms_AbvGrd               -3.441518e+03
## Full_Bath                   -3.694920e+03
## Overall_QualAverage          -5.116544e+03
## Overall_QualBelow_Average    -1.270853e+04
## Overall_QualExcellent        7.582058e+04
## Overall_QualFair             -1.147242e+04
## Overall_QualGood             1.198403e+04
## Overall_QualVery_Excellent   1.363756e+05
## Overall_QualVery_Good        3.765989e+04
## Kitchen_QualFair            -2.369556e+04
## Kitchen_QualGood            -1.611428e+04
```

```
## Kitchen_QualTypical      -2.416472e+04
## Fireplaces               1.083070e+04
## Fireplace_QuFair        -7.857118e+03
## Fireplace_QuGood         1.490291e+02
## Fireplace_QuNo_Fireplace 1.823280e+03
## Fireplace_QuPoor        -5.803365e+03
## Fireplace_QuTypical     -6.962039e+03
## Exter_QualFair          -3.297575e+04
## Exter_QualGood          -1.457261e+04
## Exter_QualTypical       -1.916633e+04
## Lot_Frontage            1.001768e+02
## Lot_Area                 6.032398e-01
## Longitude               -3.515129e+04
## Latitude                 5.776171e+04
## Misc_Val                 8.685961e-01
## Year_Sold                -5.749899e+02
```

Best tuning parameters alpha and lambda

```
enet.fit$bestTune
```

```
##      alpha  lambda
## 176  0.05 580.3529
```

test error

```
enet.pred <- predict(enet.fit, newdata = testing)
mean((enet.pred - testing[, "Sale_Price"])^2 |> pull()) #MSE
```

```
## [1] 438615743
```

```
mean((enet.pred - testing[, "Sale_Price"])^2 |> pull()) |> sqrt() #RMSE
```

```
## [1] 20943.16
```

Mechanically speaking, it is possible to apply the 1SE rule to elastic net:

```
# fit 1se version
ctrl2 <- trainControl(method = "cv", number = 10, selectionFunction = "oneSE")
set.seed(2025)
enet.1se <- train(Sale_Price ~ .,
                  data = training,
                  method = "glmnet",
                  tuneGrid = expand.grid(alpha = seq(0, 1, length = 21),
                                         lambda = exp(seq(10, -5, length = 100))),
                  trControl = ctrl2)
```

```
enet.1se$bestTune
```

```
##      alpha  lambda
## 94      0 8874.25
```

```
enet.fit$results |>
  filter(lambda == enet.fit$bestTune |> pull(lambda),
         alpha == enet.fit$bestTune |> pull(alpha)) |>
  mutate(RMSESE = RMSESD / sqrt(12),
         RMSE_1SE = RMSE + RMSESE) |>
  pull(RMSE_1SE)
```

```
## [1] 23531.45
```

```
enet.fit$results |>
  filter(RMSE > 23500,
         RMSE < 23600)
```

##	alpha	lambda	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
## 1	0.00	8874.2499	23519.64	0.8984303	16761.08	2374.337	0.01694508	1441.975
## 2	0.10	4160.2620	23519.39	0.8980426	16711.68	2355.696	0.01724977	1501.709
## 3	0.20	2640.6698	23533.60	0.8977256	16719.61	2342.145	0.01741654	1522.439
## 4	0.25	2269.4045	23551.72	0.8975295	16734.28	2345.117	0.01748456	1530.335
## 5	0.30	1950.3372	23552.90	0.8974688	16740.23	2340.162	0.01753041	1532.095
## 6	0.35	1676.1293	23540.45	0.8975203	16738.65	2327.784	0.01754849	1530.642
## 7	0.40	1440.4737	23516.82	0.8976635	16730.12	2308.848	0.01753529	1525.858
## 8	0.45	1440.4737	23581.69	0.8971593	16769.62	2334.959	0.01765471	1540.065
## 9	0.50	1237.9501	23544.89	0.8974051	16756.52	2309.808	0.01762218	1531.465
## 10	0.55	1063.9003	23500.92	0.8977170	16736.79	2283.527	0.01757529	1521.637
## 11	0.55	1237.9501	23586.43	0.8970913	16779.15	2330.421	0.01770803	1541.007
## 12	0.60	1063.9003	23548.71	0.8973405	16765.87	2301.477	0.01767086	1531.244
## 13	0.65	914.3211	23500.88	0.8976889	16744.02	2275.420	0.01760253	1521.442
## 14	0.65	1063.9003	23579.91	0.8971107	16782.52	2319.611	0.01773336	1538.715
## 15	0.70	914.3211	23539.90	0.8973802	16767.27	2290.061	0.01768902	1528.191
## 16	0.75	914.3211	23566.41	0.8971865	16781.60	2305.820	0.01773699	1535.181
## 17	0.80	785.7720	23524.48	0.8974788	16764.19	2275.506	0.01766438	1522.893
## 18	0.80	914.3211	23592.75	0.8969926	16794.16	2320.076	0.01778222	1541.487
## 19	0.85	785.7720	23547.45	0.8973089	16775.94	2289.412	0.01771678	1529.466
## 20	0.90	785.7720	23568.71	0.8971532	16786.94	2301.596	0.01775285	1534.844
## 21	0.95	675.2963	23525.42	0.8974557	16769.99	2270.027	0.01768056	1521.366
## 22	0.95	785.7720	23591.16	0.8969892	16798.55	2314.198	0.01778728	1539.736
## 23	1.00	675.2963	23542.89	0.8973281	16778.67	2281.745	0.01771693	1527.013

Although it is possible to select a 1SE lambda value using the available caret functions, I do not think this is the best idea for elastic net. It is very challenging to accurately apply the 1SE rule to the elastic net model, because there are multiple combinations of alpha and lambda that yield similar RMSE values which are close to the 1SE value. You could further refine the grid to create a more precise search for the exact 1SE value, but ultimately multiple combinations of alpha and lambda will remain as potential solutions. The new value of lambda that corresponds to the “1SE solution” also comes with a different alpha value, meaning you are then jumping onto a different curve to compare RMSE values, and thereby changing the meaning of the lambda value too. It is therefore hard to interpret which combination of values is best with this added flexibility. Thus I would not use the 1SE method with elastic net.

c. Partial least squares

- (c) Fit a partial least squares model on the training data and report the test error. How many components are included in your model?

```
set.seed(2025)

pls.fit <- train(Sale_Price ~ .,
  data = training,
  method = "pls",
  tuneGrid = data.frame(ncomp = 1:(ncol(training)-1)),
  trControl = ctrl1,
  preProcess = c("center", "scale"))
```

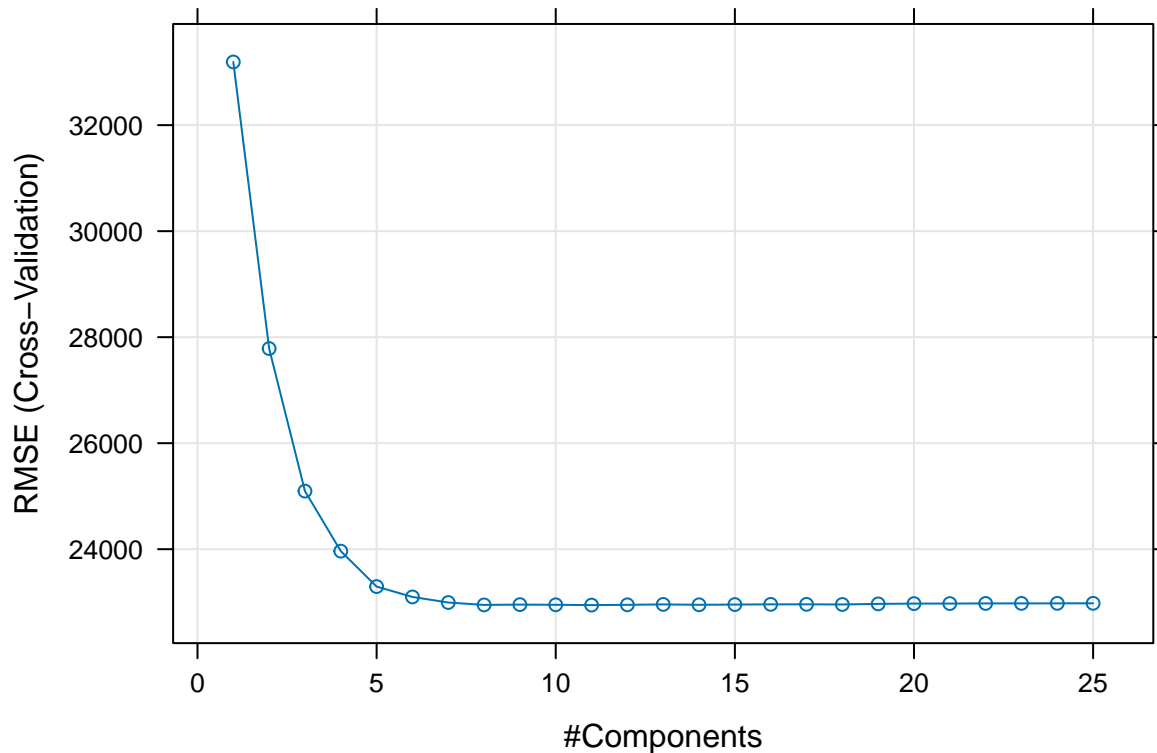
PLS model test error


```
pred_pls <- predict(pls.fit, newdata = testing)
mean((testing$Sale_Price - pred_pls)^2) |> sqrt() #RMSE
```

```
## [1] 21243.27
```

PLS number of components

```
plot(pls.fit)
```



```
pls.fit$bestTune
```

```
##      ncomp
## 11      11
```

There are 11 components in the optimal PLS model.

d. Model comparison


(d) Choose the best model for predicting the response and explain your choice.

```
resamp <- resamples(list(lasso = lasso.fit,
                        lasso_1se = lasso.1se,
                        enet = enet.fit,
                        pls = pls.fit))
summary(resamp)
```

```
##
## Call:
## summary.resamples(object = resamp)
##
## Models: lasso, lasso_1se, enet, pls
## Number of resamples: 10
```

```
##
## MAE
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## lasso      14741.86 15555.86 16284.84 16677.58 17493.56 19166.12    0
## lasso_1se  14996.43 15493.21 16421.93 16778.67 17939.52 19629.84    0
## enet       14789.53 15501.17 16268.01 16654.09 17494.45 19151.88    0
## pls       14848.76 15656.74 16411.28 16743.46 17517.21 19194.05    0
##
## RMSE
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## lasso      19793.26 21879.30 22650.75 22944.36 23782.53 27027.81    0
## lasso_1se  20825.73 21757.54 23517.66 23542.89 24701.04 28300.57    0
## enet       19827.17 21904.35 22614.81 22939.63 23774.73 27064.17    0
## pls       19837.01 21953.53 22716.43 22945.08 23772.46 26935.47    0
##
## Rsquared
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## lasso      0.8774739 0.8893619 0.9002486 0.9021463 0.9163527 0.9249577    0
## lasso_1se  0.8771573 0.8839686 0.8924394 0.8973281 0.9072272 0.9260259    0
## enet       0.8776374 0.8894404 0.9005349 0.9021885 0.9160495 0.9254490    0
## pls       0.8774820 0.8893300 0.8997448 0.9020588 0.9164506 0.9243377    0
```

`bwplot(resamp, metric = "RMSE")`



The models are fairly close, but I would opt to use the elastic net model to predict the Sale Price. It is a fairly flexible model and here yields the lowest average RMSE. The lasso and PLS models are both quite close in performance as well. The lasso 1SE model has worse prediction performance, but has fewer predictors and so could be preferred on that basis. However in this case, when I am just trying to optimize predictive power, I would choose the elastic net model.

That choice is made solely using training error. At this point we can compare the models based on test error.

We can NOT use these results to inform model selection (since we specifically set aside the testing data for this purpose), but it can be interesting to see and understand how our chosen model did on out-of-sample data.

Testing error for each model

```
# Lasso
mean((lasso.pred - testing$Sale_Price)^2) |> sqrt()
```

```
## [1] 20965.89
```

```
# Lasso 1SE
mean((predict(lasso.1se,newdata = testing) - testing$Sale_Price)^2) |> sqrt()
```

```
## [1] 20492.55
```

```
# Elastic Net
mean((enet.pred - testing$Sale_Price)^2) |> sqrt()
```

```
## [1] 20943.16
```

```
# PLS
mean((pred_pls - testing$Sale_Price)^2) |> sqrt()
```

```
## [1] 21243.27
```

It turned out that Elastic Net had the 2nd lowest test error, with the lowest belonging to the Lasso 1SE. Personally this is a good lesson for me for the future when it comes to model selection. The 1SE rule consciously selects a model that has worse training error than other “best fit” models. However by selecting a parsimonious set of predictors it may perform better on out of sample data. In all, I am satisfied with my choice of the elastic net model, since the testing error was fairly close to my expectations based on the training error.

e. Lasso using glmnet

- (e) If R package “caret” was used for the lasso in (a), retrain this model using R package “glmnet”, and vice versa. Compare the selected tuning parameters between the two software approaches. Should there be discrepancies in the chosen parameters, discuss potential reasons for these differences.


Fitting lasso using glmnet

```
x <- model.matrix(Sale_Price ~ ., training)[,-1]
y <- training$Sale_Price

set.seed(2025)

lasso.glmnet <- cv.glmnet(x, y,
                          alpha = 1,
                          lambda = exp(seq(10, -5, length = 100)))

plot(lasso.glmnet)
```



```
lasso.glmnet$lambda.min
```

```
## [1] 69.57632
```

```
lasso.glmnet$lambda.1se
```

```
## [1] 785.772
```

```
# compare to the values from caret
```

```
lasso.fit$bestTune
```

```
##      alpha      lambda
```

```
## 62      1 69.57632
```

```
lasso.1se$bestTune
```

```
##      alpha      lambda
```

```
## 77      1 675.2963
```

The minimum lambda value using glmnet and caret was exactly the same. The 1SE lambda value was slightly different (~785 for glmnet and ~675 for caret), but a plausible explanation for this may caret choosing the first value below the 1SE cutoff and glmnet choosing the first value above above the 1SE cutoff, rather than a more substantial estimation difference between methods.